

CSCI 180

Buffer Overflow Exercise

1. Overview

The learning objective of this lab is for you to gain the first-hand experience on buffer-overflow vulnerability. Buffer overflow is defined as the condition in which a program attempts to write data beyond the boundaries of pre-allocated fixed length buffers. This vulnerability can be used by a malicious user to alter the flow control of the program, leading to the execution of malicious code. This vulnerability arises due to the mixing of the storage for data (e.g. buffers) and the storage for controls (e.g. return addresses): an overflow in the data part can affect the control flow of the program, because an overflow can change the return address.

Lab environment: Use the pre-built Ubuntu VM that has been provided for this class. **Note:** If you are using a shared folder between your host machine and your VM, make sure to not run the codes from inside the shared folder. Copy them to some other location on Ubuntu and run from there.

Submission: You need to submit a detailed lab report to describe what you have done and what you have observed. Follow the tasks and for each task answer the Q# specifically in your report. You may provide explanation to the observations that are interesting or surprising. You can always add code snippets or screenshots of what you have observed. You are encouraged to pursue further investigation, beyond what is required by the lab description. You can earn bonus points for extra efforts (at the discretion of your instructor). Only submit typed reports electronically. No handwritten reports accepted.

2. Tasks

Task 1: Setting up the environment by disabling countermeasures

As previously discussed in class, Operating systems and compilers have implemented several security mechanisms to make the buffer overflow attack more difficult. In this section, in order to make the attack easier, we will disable these countermeasures.

- **Address Space Randomization.** In order to disable address randomization temporarily use the following command:

```
$ sudo sysctl -w kernel.randomize_va_space=0
```

Remember that the changes we make are temporary, so if you reboot your system you will need to run this command again. To view the current settings for ASLR on your system you can use this command:

```
$ sysctl -a --pattern randomize
```

If the value is 2, it means full randomization is on, and if it is 0 it means that randomization is disabled.

- **StackGuard protection** scheme is a method used by GCC compiler to prevent buffer overflow. We will disable this feature by using the `-fno-stack-protector` option while compiling the program.
- **Nonexecutable stack.** The most recent version of gcc automatically marks the binary version of program to indicate that this program does not require executable stack. In order to disable this countermeasure we use the option `-z execstack` when compiling the program to make the stack executable.

- **Configuring /bin/sh** As discussed in Lab1, the shell in Ubuntu 16.04 VM has a countermeasure that prevents itself from being executed in a Set-UID process. In this lab, our victim program is a Set-UID program, and our attack is trying to run /bin/sh, therefore this countermeasure makes our attack more difficult. We will make our attack easier by linking /bin/sh to another shell that does not have such a countermeasure. Later we will see that it is not very hard for the attacker to defeat this countermeasure as part of the attack.

```
$ sudo rm /bin/sh
$ sudo ln -s /bin/zsh /bin/sh
```

Task 2: The shellcode

Before doing the attack, let's get familiar with the shellcode. A shellcode is a code that launches a shell. This is usually the malicious code injected by the attacker to the stack. A program called `shellcodetest.c` is given to you. Take a look at this file. `code[]` contains the program to execute a shell in bytecode. It is as if you write a program to execute a shell in C and then compile and extract the byte code from it.

`shellcodetest.c` program will test this bytecode to make sure it will launch a shell. Compile this program using the following gcc command.

```
$ gcc -z execstack -o shellcodetest shellcodetest.c
```

Q1: Run the program and describe your observations.

Task 3: The vulnerable program

You are provided with a file called `unsafe.c`. We are assuming this is a root-owned Set-UID program that has a buffer overflow vulnerability in it. Your goal in this lab is to exploit this vulnerability to spawn a shell with root privilege.

Compile this vulnerable program by including the `-fno-stack-protector` and `-z execstack` options to turn off the StackGuard and the non-executable stack protections. After the compilation, we need to make the program a root-owned Set-UID program.

```
$ gcc -z execstack -fno-stack-protector -o unsafe unsafe.c
$ sudo chown root unsafe
$ sudo chmod 4755 unsafe
```

This program has a buffer overflow vulnerability as indicated in the code. It first reads an input from a file called `inputfile`, and then passes this input to another buffer in the function `copyfunc()`. The original input can have a maximum length of 517 bytes, but the buffer in `copyfunc()` is only 24 bytes long. Because `strcpy()` does not check boundaries, buffer overflow will occur. Since this program is a root-owned Set-UID program, if a normal user can exploit this buffer overflow vulnerability, the normal user might be able to get a shell with root privileges. Remember that this program receives its input from a file and this file is under users' control. Our objective is to create the contents for `inputfile`, such that when the vulnerable program copies the contents into its buffer, a root shell can be spawned.

Task 4: Exploiting the vulnerability, the real attack

You are given a partially completed exploit code called `exploit.py`. This code is written in Python for simplicity. The goal of this code is to construct contents for `inputfile`. In this code, the shellcode is given to you. You need to develop the rest.

Your task is to (1) finish this program, (2) run it by running `./exploit.py`. This will generate the contents for `inputfile` in a way that when the vulnerable program `unsafe` reads this file and copies its contents into the stack, the buffer overflow vulnerability results in executing this shellcode and invoking a shell.

troubleshoot: If you receive a permission denied error when running `exploit.py`, you may need to give it execute

permission: `chmod a+x exploit.py`

After you created the inputfile by running the `exploit.py`, (3) run the vulnerable program `unsafe`. If your exploit is implemented correctly, you should be able to get a root shell with `#` sign. Once you receive the shell you can check your effective uid and real uid by typing `id` on the screen.

Q2: Explain your process by showing your `exploit.py` code or showing the value of `D` and the value you replaced for `content[D+0 to 3]`. Explain how you came up with these values. Show a screenshot of the result when you run `unsafe` program, and the result of `id` from the shell you invoked.

Help with gdb

You can use `gdb` to get some more information about the vulnerable program. Compile the code to debug the program:

```
$ gcc unsafe.c -o unsafe_gdb -g -z execstack -fno-stack-protector
```

Create an empty inputfile

```
$ touch inputfile
```

Run the debugger

```
$ gdb unsafe_gdb
```

Inside the debugger, create a breakpoint at the `copyfunc()`

```
gdb-peda$ b copyfunc
```

You can print the address of the `buffer` using the command:

```
gdb-peda$ p &buffer
```

You can print the `ebp` address using the command:

```
gdb-peda$ p $ebp
```

Use `Control+D` to exit from GDB at any time.

Task 5: Defeating shell's countermeasure

As explained before, the shell in Ubuntu 16.04 drops privileges when it detects that the effective UID does not equal to the real UID. This is what we saw in Lab1, and in order to make our attack work, we changed the shell we are using in Ubuntu to a version that does not have this countermeasure. In this task, we want to see that even if we don't disable this countermeasure, the attacker is able to defeat this. In order to be able to test this, you first need to change the shell back to a version with the countermeasure by using the following commands:

```
$ sudo rm /bin/sh
$ sudo ln -s /bin/dash /bin/sh
```

One approach for the attacker to defeat this countermeasure is to change the real user ID of the victim process to zero before invoking the shell program. This way, when the shell compares the effective uid and the real uid they both

will be equal to zero. This is because the effective uid is 0 as a result of the program being a root-owned Set-UID program. We can achieve this by calling `setuid(0)` before invoking `/bin/sh` in the shellcode. The bytecode for calling `setuid(0)` is as follows. All you need to do is to add the following code to the beginning of the shellcode in `exploit.py`:

```
"\x31\xc0"      /* xorl    %eax,%eax      */
"\x31\xdb"      /* xorl    %ebx,%ebx      */
"\xb0\xd5"      /* movb    $0xd5,%al      */
"\xcd\x80"      /* int     $0x80           */
```

Q3: Try the attack from Task 4 again and see if you can get a root shell. Check the uid by typing `id` and report it. Please describe your results.

Task 6: Defeating Address Randomization

As mentioned in class, on 32-bit Linux machines, stacks randomization space is not very high and can be brute forced. In this task we will try to defeat the address randomization countermeasure by brute forcing. First we need to turn on the Ubuntu's address randomization using the following command.

```
$ sudo /sbin/sysctl -w kernel.randomize_va_space=2
```

Q4: Run the same attack as in Task 4 and report your observation.

We now try to attack the vulnerable program repeatedly, hoping that the address we put in the `inputfile` can eventually be correct. You can use the following shell script to run the vulnerable program in an infinite loop. If the attack succeeds, the script will stop. This may take a while, try it as long as it is possible depending on your system. You might want to let it run overnight if needed.

```
#!/bin/bash

SECONDS=0
value=0
while [ 1 ]
do
    value=$(( $value+1 ))
    duration=$SECONDS
    min=$(( $duration / 60 ))
    sec=$(( $duration % 60 ))
    echo "The program has been running $value times so far."
    echo "It has been $min:$sec (mins:secs)"
    ./unsafe
    echo ""
done
```

Q5: Run this code and describe your observations. If you did not get to successfully run the attack after a reasonable time, stop the program. Show a screenshot of how long you ran the script.

Note for those who are not familiar with writing shell scripts: Write the provided code in a file, suppose you name it `myattack`. Then run this file: `./myattack`. If you get a permission denied error, give the file execution permission by using: `chmod a+x myattack` and run it again.